

# Machine Learning with SVM: House Price Prediction & Digit Classification

## SVM (support vector machine)

- Regression problem
- classification problem
  - Binary
  - multiclass

In [ ]: *# Designed By : ALTAF HUSAIN DATA ANALYST*

## Housing Price Prediction : Regression Problem (SVR)



## step 1 : modules import

```
In [80]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC, SVR
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import confusion_matrix, classification_report, mean_absolute_error, mean_squared_error, root_mean_squared_error
from sklearn.datasets import load_digits, fetch_california_housing
import warnings
```

```
warnings.filterwarnings('ignore')
print("All modules loaded succesfully")
```

All modules loaded succesfully

```
In [81]: house = fetch_california_housing()
df = pd.DataFrame(house.data, columns= house.feature_names)
df[house.target_names[0]] = house.target
```

```
In [82]: df.sample()
```

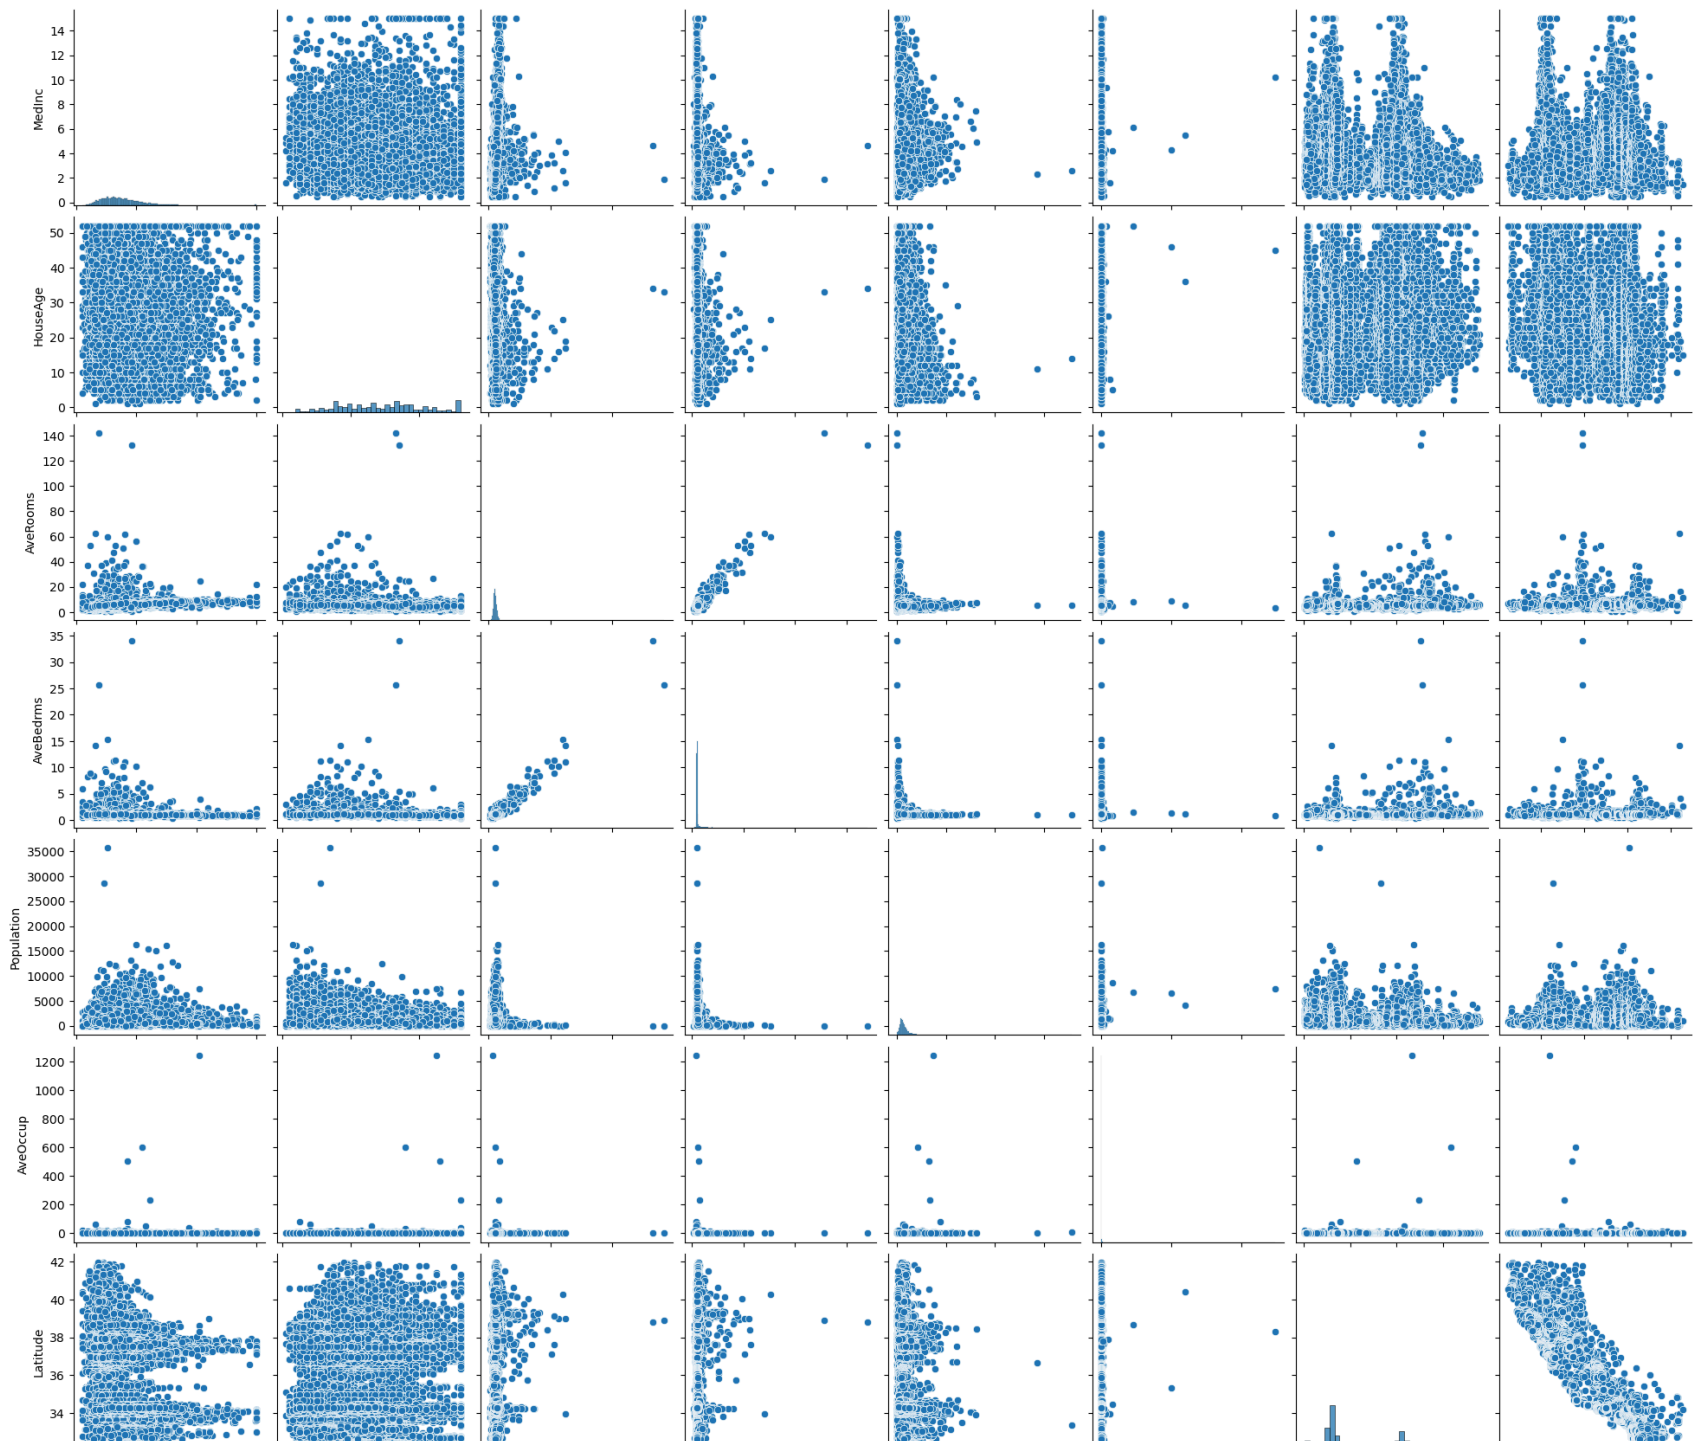
```
Out[82]:
```


	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MedHouseVal
<b>13444</b>	2.1189	13.0	4.47792	1.164868	2475.0	2.428852	34.08	-117.43	1.217

```
In [83]: # print(house['DESCR'])
```

```
In [84]: X = df.iloc[:, :-1]
y = df['MedHouseVal']
```

```
In [85]: sns.pairplot(X)
plt.show()
```





```
In [86]: X.sample()
```

```
Out[86]:
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
<b>853</b>	4.6786	16.0	5.350061	1.088127	2530.0	3.096695	37.59	-122.03

## MIN\_MAX SCALER

```
In [87]: scaler = MinMaxScaler() # scale to 0 and 1
```

```
In [88]: X_scaled = scaler.fit_transform(X)
```

```
In [89]: X_scaled.describe()
```

Out[89]:

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
<b>count</b>	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
<b>mean</b>	3.870671	28.639486	5.429000	1.096675	1425.476744	3.070655	35.631861	-119.569704
<b>std</b>	1.899822	12.585558	2.474173	0.473911	1132.462122	10.386050	2.135952	2.003532
<b>min</b>	0.499900	1.000000	0.846154	0.333333	3.000000	0.692308	32.540000	-124.350000
<b>25%</b>	2.563400	18.000000	4.440716	1.006079	787.000000	2.429741	33.930000	-121.800000
<b>50%</b>	3.534800	29.000000	5.229129	1.048780	1166.000000	2.818116	34.260000	-118.490000
<b>75%</b>	4.743250	37.000000	6.052381	1.099526	1725.000000	3.282261	37.710000	-118.010000
<b>max</b>	15.000100	52.000000	141.909091	34.066667	35682.000000	1243.333333	41.950000	-114.310000

In [90]: `pd.DataFrame(X_scaled).describe()`

Out[90]:

	0	1	2	3	4	5	6	7
<b>count</b>	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000	20640.000000
<b>mean</b>	0.232464	0.541951	0.032488	0.022629	0.039869	0.001914	0.328572	0.476125
<b>std</b>	0.131020	0.246776	0.017539	0.014049	0.031740	0.008358	0.226988	0.199555
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	0.142308	0.333333	0.025482	0.019943	0.021974	0.001398	0.147715	0.253984
<b>50%</b>	0.209301	0.549020	0.031071	0.021209	0.032596	0.001711	0.182784	0.583665
<b>75%</b>	0.292641	0.705882	0.036907	0.022713	0.048264	0.002084	0.549416	0.631474
<b>max</b>	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

## Train-test part

```
In [91]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size = 0.2 ,random_state=42)
```

```
In [92]: X_train.shape
```

```
Out[92]: (16512, 8)
```

```
In [93]: X_test.shape
```

```
Out[93]: (4128, 8)
```

```
In [94]: y_train.shape
```

```
Out[94]: (16512,)
```

```
In [95]: y_test.shape
```


```
Out[95]: (4128,)
```

## Model Building

```
In [96]: model_reg = SVR()
```

```
In [97]: model_reg.fit(X_train,y_train)
```

```
Out[97]:
```



```
In [98]: y_pred = model_reg.predict(X_test)
```

```
In [99]: mae = mean_absolute_error(y_test,y_pred)
print("Mae : ",mae)
```

```
Mae : 0.8599506583445776
```

```
In [100]: mse = mean_squared_error(y_test,y_pred)
print("Mse : ",mse)
```

Mse : 1.3320115421348788

```
In [101]: rmse = root_mean_squared_error(y_test,y_pred)
print("RMse : ",rmse)
```

RMse : 1.1541280440812791

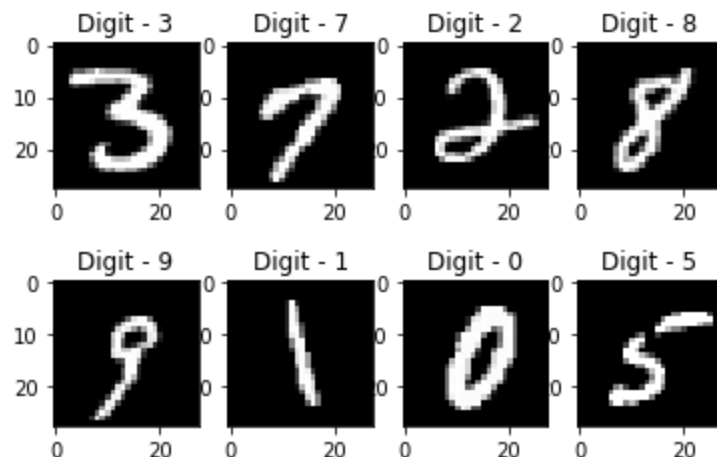
```
In [102]: model_reg.score(X_test,y_test)
```

Out[102]: -0.01648536010717705

```
In [103]: sample_data = pd.DataFrame(X_scaled,columns = house.feature_names).sample()
price = round(model_reg.predict(sample_data)[0]* 100000,2)
print(f'MedHouseVal : ${price}')
```

MedHouseVal : \$140943.73

## Digits classification : multiclass Classification problem (SVC)



```
In [104]: digits = load_digits()
```



```
In [105]: df = pd.DataFrame(digits.data, columns= digits.feature_names)
df['target'] = digits.target
```

```
In [106]: df.sample()
```

```
Out[106]:
```

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_7	pixel
1222	0.0	0.0	0.0	9.0	9.0	0.0	0.0	0.0	0.0	0.0	...	1.0	

1 rows × 65 columns




```
In [107]: df.shape
```

```
Out[107]: (1797, 65)
```

```
In [108]: X = df.iloc[:, :-1]
```

```
In [109]: y = df['target']
```

```
In [110]: image_digit = X.sample().to_numpy().reshape(8,8)
plt.figure(figsize=(1,1))
plt.gray()
plt.imshow(image_digit)
plt.show()
```



```
In [111]: df['target'].unique()
```

```
Out[111]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [112]: image_digit
```

```
Out[112]: array([[ 0.,  1., 13., 16., 16., 16., 12.,  1.],
 [ 0.,  6., 16., 14., 12., 11.,  5.,  0.],
 [ 0.,  2., 15., 15.,  5.,  0.,  0.,  0.],
 [ 0.,  0.,  8., 14., 15.,  1.,  0.,  0.],
 [ 0.,  0.,  0.,  3., 16.,  6.,  0.,  0.],
 [ 0.,  0.,  0.,  3., 16.,  5.,  0.,  0.],
 [ 0.,  0.,  7., 10., 16.,  4.,  0.,  0.],
 [ 0.,  0., 15., 16., 10.,  0.,  0.,  0.]])
```

```
In [113]: # image : 0 - 255
```

```
In [114]: scaler = MinMaxScaler()
```

```
In [115]: X_scaled = scaler.fit_transform(X)
```

```
In [116]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)
```

```
In [117]: print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(1437, 64)
```

```
(360, 64)
```

```
(1437,)
```

```
(360,)
```

```
In [118]: model_cl = SVC()
```

```
In [119]: model_cl.fit(X_train,y_train)
```

```
Out[119]: SVC
SVC()
```


```
In [120]: y_pred = model_cl.predict(X_test)
```

```
In [121]: temp_df = pd.DataFrame({'Actual class':y_test,'Predicted class':y_pred})
```

```
In [122]: cm =confusion_matrix(y_test,y_pred)
cm
```

```
Out[122]: array([[33,  0,  0,  0,  0,  0,  0,  0,  0,  0],
 [ 0, 28,  0,  0,  0,  0,  0,  0,  0,  0],
 [ 0,  0, 33,  0,  0,  0,  0,  0,  0,  0],
 [ 0,  0,  0, 33,  0,  1,  0,  0,  0,  0],
 [ 0,  0,  0,  0, 46,  0,  0,  0,  0,  0],
 [ 0,  0,  0,  0,  0, 46,  1,  0,  0,  0],
 [ 0,  0,  0,  0,  0,  0, 35,  0,  0,  0],
 [ 0,  0,  0,  0,  0,  0,  0, 33,  0,  1],
 [ 0,  0,  0,  0,  0,  0,  0,  0, 29,  1],
 [ 0,  0,  0,  0,  0,  1,  0,  1,  0, 38]])
```

```
In [123]: sns.heatmap(cm,annot =True)
plt.xlabel('Predicted result')
plt.ylabel('Actual result')
plt.show()
```



```
In [124]: model_score = model_cl.score(X_test,y_test)
```

```
In [125]: print(f'Model score is :{round(model_score *100,2)}%')
```

Model score is :98.33%


```
In [126]: print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	33
1	1.00	1.00	1.00	28
2	1.00	1.00	1.00	33
3	1.00	0.97	0.99	34
4	1.00	1.00	1.00	46
5	0.96	0.98	0.97	47
6	0.97	1.00	0.99	35
7	0.97	0.97	0.97	34
8	1.00	0.97	0.98	30
9	0.95	0.95	0.95	40
accuracy			0.98	360
macro avg	0.99	0.98	0.98	360
weighted avg	0.98	0.98	0.98	360

```
In [127]: sample_data = pd.DataFrame(X_scaled).sample()
ans = model_cl.predict(sample_data)[0]

plt.figure(figsize=(1,1))
plt.gray()
plt.title(f'Predicted digit is : {ans}')
plt.imshow(sample_data.to_numpy().reshape(8,8))
plt.show()
```

Predicted digit is : 9



## Giving image to predict

```
In [130]: arr = plt.imread("C:\\Users\\SK COMPUTER\\Downloads\\img.jpg")
```


```
In [131]: img_arr = arr.copy()
```

```
In [132]: img_arr.shape
```

```
Out[132]: (8, 8, 3)
```

```
In [133]: plt.imshow(img_arr[:, :, 0])
```

```
Out[133]: <matplotlib.image.AxesImage at 0x2814e394e10>
```



```
In [134]: img_arr = img_arr[:, :, 0] if img_arr.ndim == 3 else img_arr
```

```
In [135]: from skimage.transform import resize  
img_arr = resize(img_arr, (8, 8), anti_aliasing=True)
```

```
In [136]: sample_data = np.expand_dims(img_arr.flatten(),axis = 0)
```

```
In [137]: sample_scaled_data = scaler.transform(sample_data)
```

```
In [138]: ans = model_cl.predict(sample_scaled_data)[0]  
print("Predicted digit is:", ans)
```

Predicted digit is: 4

```
In [ ]: # Designed By : ALTAF HUSAIN DATA ANALYST
```